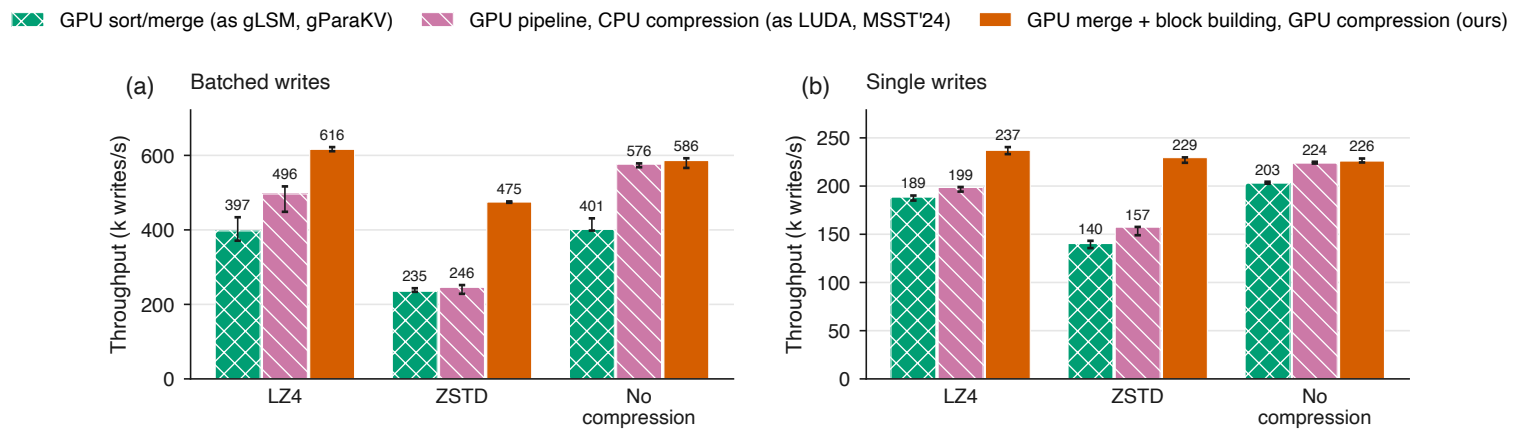
GPU sort/merge (as gLSM, gParaKV)
GPU pipeline, CPU compression (as LUDA, MSST'24)
GPU merge + block building, GPU compression (ours)
(a)
Batched writes
Throughput (k writes/s)
0
200
400
600
397
496
616
235
246
475
401
576
586
LZ4
ZSTD
No compression
(b)
Single writes
Throughput (k writes/s)
0
50
100
150
200
250
189
199
237
140
157
229
203
224
226
LZ4
ZSTD
No compression